

# Approximate Functions

# Going Deep

Jonathon Hare

Vision, Learning and Control  
University of Southampton

- No free lunch and universal approximation
- Why go deep?
- Problems of going deep
- Some fixes:
  - Improving gradient flow with skip connections
  - Regularising with Dropout

# No Free Lunch

- Statistical learning theory claims that a machine can generalise well from a finite training set.

---

<sup>1</sup>or perhaps more generally rules which are not certain

# No Free Lunch

- Statistical learning theory claims that a machine can generalise well from a finite training set.
- This contradicts basic inductive reasoning which says to derive a rule describing every member of a set one must have information about every member.

---

<sup>1</sup>or perhaps more generally rules which are not certain

# No Free Lunch

- Statistical learning theory claims that a machine can generalise well from a finite training set.
- This contradicts basic inductive reasoning which says to derive a rule describing every member of a set one must have information about every member.
- Machine learning avoids this problem by learning probabilistic<sup>1</sup> rules which are *probably* correct about *most* members of the set they concern.

---

<sup>1</sup>or perhaps more generally rules which are not certain

# No Free Lunch

- Statistical learning theory claims that a machine can generalise well from a finite training set.
- This contradicts basic inductive reasoning which says to derive a rule describing every member of a set one must have information about every member.
- Machine learning avoids this problem by learning probabilistic<sup>1</sup> rules which are *probably* correct about *most* members of the set they concern.
- But, **no free lunch theorem** states that every possible classification machine has the *same error* when averaged over *all possible* data-generating distributions.

---

<sup>1</sup>or perhaps more generally rules which are not certain

# No Free Lunch

- Statistical learning theory claims that a machine can generalise well from a finite training set.
- This contradicts basic inductive reasoning which says to derive a rule describing every member of a set one must have information about every member.
- Machine learning avoids this problem by learning probabilistic<sup>1</sup> rules which are *probably* correct about *most* members of the set they concern.
- But, **no free lunch theorem** states that every possible classification machine has the *same error* when averaged over *all possible* data-generating distributions.
  - **No machine learning algorithm is universally better than any other!**

---

<sup>1</sup>or perhaps more generally rules which are not certain



# No Free Lunch

- Statistical learning theory claims that a machine can generalise well from a finite training set.
- This contradicts basic inductive reasoning which says to derive a rule describing every member of a set one must have information about every member.
- Machine learning avoids this problem by learning probabilistic<sup>1</sup> rules which are *probably* correct about *most* members of the set they concern.
- But, **no free lunch theorem** states that every possible classification machine has the *same error* when averaged over *all possible* data-generating distributions.
  - **No machine learning algorithm is universally better than any other!**
  - Fortunately, in the real world, data is generated by a small subset of generating distributions...

---

<sup>1</sup>or perhaps more generally rules which are not certain

# The Universal Approximation Theorem

Let  $\psi : \mathbb{R} \rightarrow \mathbb{R}$  be a nonconstant, bounded, and continuous function.

# The Universal Approximation Theorem

Let  $\psi : \mathbb{R} \rightarrow \mathbb{R}$  be a nonconstant, bounded, and continuous function. Let  $I_m$  denote the  $m$ -dimensional unit hypercube  $[0, 1]^m$ .

# The Universal Approximation Theorem

Let  $\psi : \mathbb{R} \rightarrow \mathbb{R}$  be a nonconstant, bounded, and continuous function. Let  $I_m$  denote the  $m$ -dimensional unit hypercube  $[0, 1]^m$ . The space of real-valued continuous functions on  $I_m$  is denoted by  $C(I_m)$ .

# The Universal Approximation Theorem

Let  $\psi : \mathbb{R} \rightarrow \mathbb{R}$  be a nonconstant, bounded, and continuous function. Let  $I_m$  denote the  $m$ -dimensional unit hypercube  $[0, 1]^m$ . The space of real-valued continuous functions on  $I_m$  is denoted by  $C(I_m)$ . Then, given any  $\varepsilon > 0$  and any function  $f \in C(I_m)$ , there exist an integer  $N$ , real constants  $v_i, b_i \in \mathbb{R}$  and real vectors  $w_i \in \mathbb{R}^m$  for  $i = 1, \dots, N$ , such that we may define:

$F(x) = \sum_{i=1}^N v_i \psi(w_i^T x + b_i)$  as an approximate realization of the function  $f$  ; that is,

# The Universal Approximation Theorem

Let  $\psi : \mathbb{R} \rightarrow \mathbb{R}$  be a nonconstant, bounded, and continuous function. Let  $I_m$  denote the  $m$ -dimensional unit hypercube  $[0, 1]^m$ . The space of real-valued continuous functions on  $I_m$  is denoted by  $C(I_m)$ . Then, given any  $\varepsilon > 0$  and any function  $f \in C(I_m)$ , there exist an integer  $N$ , real constants  $v_i, b_i \in \mathbb{R}$  and real vectors  $w_i \in \mathbb{R}^m$  for  $i = 1, \dots, N$ , such that we may define:

$F(x) = \sum_{i=1}^N v_i \psi(w_i^T x + b_i)$  as an approximate realization of the function  $f$  ; that is,

$$|F(x) - f(x)| < \varepsilon \quad \forall \quad x \in I_m.$$

# The Universal Approximation Theorem

Let  $\psi : \mathbb{R} \rightarrow \mathbb{R}$  be a nonconstant, bounded, and continuous function. Let  $I_m$  denote the  $m$ -dimensional unit hypercube  $[0, 1]^m$ . The space of real-valued continuous functions on  $I_m$  is denoted by  $C(I_m)$ . Then, given any  $\varepsilon > 0$  and any function  $f \in C(I_m)$ , there exist an integer  $N$ , real constants  $v_i, b_i \in \mathbb{R}$  and real vectors  $w_i \in \mathbb{R}^m$  for  $i = 1, \dots, N$ , such that we may define:

$F(x) = \sum_{i=1}^N v_i \psi(w_i^T x + b_i)$  as an approximate realization of the function  $f$  ; that is,

$$|F(x) - f(x)| < \varepsilon \quad \forall \quad x \in I_m.$$

$\implies$  simple neural networks can represent a wide variety of interesting functions when given appropriate parameters.

# So a single hidden layer network can approximate most functions?

- Yes!

---

<sup>2</sup>note that it has been shown that the gradients of the function are approximated by the network to an arbitrary precision



# So a single hidden layer network can approximate most functions?

- Yes!
- But, ...
  - worse-case analysis shows it might be exponential (possibly one hidden unit for *every* input configuration)

---

<sup>2</sup>note that it has been shown that the gradients of the function are approximated by the network to an arbitrary precision

# So a single hidden layer network can approximate most functions?

- Yes!
- But, ...
  - to get the precision you require (small  $\varepsilon$ ), you might need a really large number of hidden units (very large  $N$ ).
  - worse-case analysis shows it might be exponential (possibly one hidden unit for *every* input configuration)

---

<sup>2</sup>note that it has been shown that the gradients of the function are approximated by the network to an arbitrary precision

# So a single hidden layer network can approximate most functions?

- Yes!
- But, ...
  - to get the precision you require (small  $\varepsilon$ ), you might need a really large number of hidden units (very large  $N$ ).
  - worse-case analysis shows it might be exponential (possibly one hidden unit for every input configuration)
  - We've not said anything about learnability...

---

<sup>2</sup>note that it has been shown that the gradients of the function are approximated by the network to an arbitrary precision

# So a single hidden layer network can approximate most functions?

- Yes!
- But, ...
  - to get the precision you require (small  $\varepsilon$ ), you might need a really large number of hidden units (very large  $N$ ).
  - worse-case analysis shows it might be exponential (possibly one hidden unit for every input configuration)
  - We've not said anything about learnability...
    - The optimiser might not find a good solution<sup>2</sup>.

---

<sup>2</sup>note that it has been shown that the gradients of the function are approximated by the network to an arbitrary precision

# So a single hidden layer network can approximate most functions?

- Yes!
- But, ...
  - to get the precision you require (small  $\varepsilon$ ), you might need a really large number of hidden units (very large  $N$ ).
  - worse-case analysis shows it might be exponential (possibly one hidden unit for every input configuration)
  - We've not said anything about learnability...
    - The optimiser might not find a good solution<sup>2</sup>.
    - The training algorithm might just choose the wrong solution as a result of overfitting.

---

<sup>2</sup>note that it has been shown that the gradients of the function are approximated by the network to an arbitrary precision

# So a single hidden layer network can approximate most functions?

- Yes!
- But, ...
  - to get the precision you require (small  $\varepsilon$ ), you might need a really large number of hidden units (very large  $N$ ).
  - worse-case analysis shows it might be exponential (possibly one hidden unit for every input configuration)
  - We've not said anything about learnability...
    - The optimiser might not find a good solution<sup>2</sup>.
    - The training algorithm might just choose the wrong solution as a result of overfitting.
    - *There is no known universal procedure for examining a set of examples and choosing a function that will generalise to points out of the training set.*

---

<sup>2</sup>note that it has been shown that the gradients of the function are approximated by the network to an arbitrary precision

# Then Why Go Deep?

- There are functions you can compute with a deep neural network that shallow networks require exponentially more hidden units to compute.
- The following function is more efficient to implement using a deep neural network:  $y = x_1 \oplus x_2 \oplus x_3 \oplus \cdots \oplus x_n$

# Vanishing and Exploding Gradients

- The vanishing and exploding gradient problem is a difficulty found in training NN with gradient-based learning methods and backpropagation.



# Vanishing and Exploding Gradients

- The vanishing and exploding gradient problem is a difficulty found in training NN with gradient-based learning methods and backpropagation.
- In training, the gradient may become vanishingly small (or large), effectively preventing the weight from changing its value (or exploding in value).

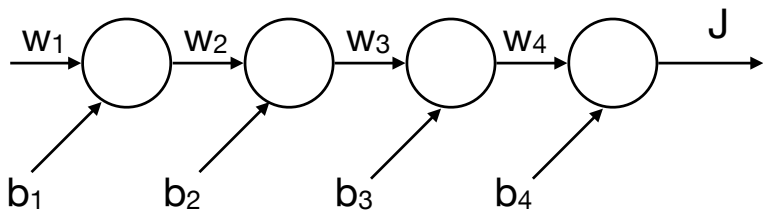
# Vanishing and Exploding Gradients

- The vanishing and exploding gradient problem is a difficulty found in training NN with gradient-based learning methods and backpropagation.
- In training, the gradient may become vanishingly small (or large), effectively preventing the weight from changing its value (or exploding in value).
- This leads to the neural network not being able to train.

# Vanishing and Exploding Gradients

- The vanishing and exploding gradient problem is a difficulty found in training NN with gradient-based learning methods and backpropagation.
- In training, the gradient may become vanishingly small (or large), effectively preventing the weight from changing its value (or exploding in value).
- This leads to the neural network not being able to train.
- This issue affects many-layered networks (feed-forward), as well as recurrent networks.

# Issues with Going Deep



# Residual Connections

- One of the most effective ways to resolve the vanishing gradient problem is with residual neural networks (ResNets)<sup>3</sup>.

---

<sup>3</sup>K. He, X. Zhang, S. Ren and J. Sun, "Deep Residual Learning for Image Recognition," CVPR, Las Vegas, NV, 2016, pp. 770-778.

# Residual Connections

- One of the most effective ways to resolve the vanishing gradient problem is with residual neural networks (ResNets)<sup>3</sup>.
- ResNets are artificial neural networks that use *skip connections* to jump over layers.

---

<sup>3</sup>K. He, X. Zhang, S. Ren and J. Sun, "Deep Residual Learning for Image Recognition," CVPR, Las Vegas, NV, 2016, pp. 770-778.

# Residual Connections

- One of the most effective ways to resolve the vanishing gradient problem is with residual neural networks (ResNets)<sup>3</sup>.
- ResNets are artificial neural networks that use *skip connections* to jump over layers.
- The vanishing gradient problem is mitigated in ResNets by reusing activations from a previous layer.

---

<sup>3</sup>K. He, X. Zhang, S. Ren and J. Sun, "Deep Residual Learning for Image Recognition," CVPR, Las Vegas, NV, 2016, pp. 770-778.

# Residual Connections

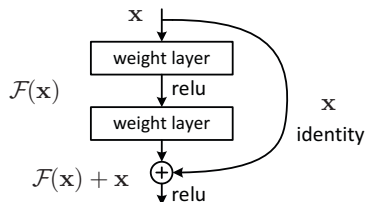


Figure 2. Residual learning: a building block.<sup>4</sup>

<sup>4</sup>K. He, X. Zhang, S. Ren and J. Sun, "Deep Residual Learning for Image Recognition," CVPR, Las Vegas, NV, 2016, pp. 770-778.



# Residual Connections

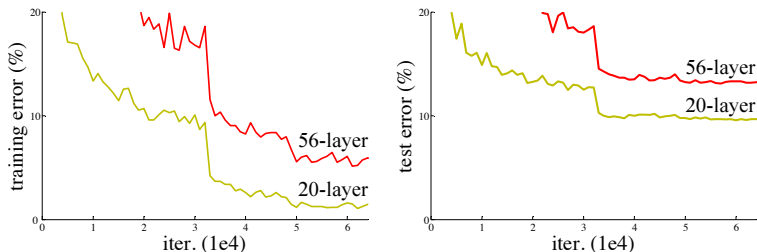


Figure 1. Training error (left) and test error (right) on CIFAR-10 with 20-layer and 56-layer “plain” networks. The deeper network has higher training error, and thus test error. Similar phenomena on ImageNet is presented in Fig. 4.

5

<sup>5</sup>K. He, X. Zhang, S. Ren and J. Sun, “Deep Residual Learning for Image Recognition,” CVPR, Las Vegas, NV, 2016, pp. 770-778.

# Residual Connections

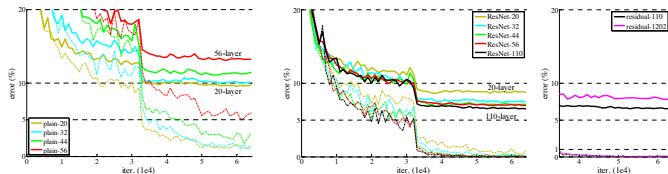


Figure 6. Training on **CIFAR-10**. Dashed lines denote training error, and bold lines denote testing error. **Left:** plain networks. The error of plain-110 is higher than 60% and not displayed. **Middle:** ResNets. **Right:** ResNets with 110 and 1202 layers.

6

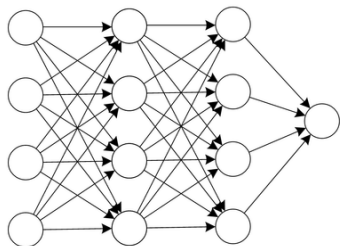
<sup>6</sup>K. He, X. Zhang, S. Ren and J. Sun, "Deep Residual Learning for Image Recognition," CVPR, Las Vegas, NV, 2016, pp. 770-778.

- Neural networks with a large number of parameters (and hidden layers) are powerful, however, overfitting is a serious problem in such systems.

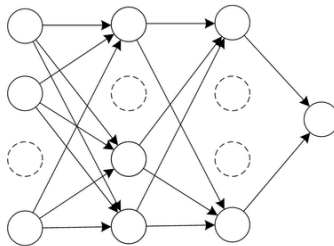
- Neural networks with a large number of parameters (and hidden layers) are powerful, however, overfitting is a serious problem in such systems.
- Dropout is a form of regularization

- Neural networks with a large number of parameters (and hidden layers) are powerful, however, overfitting is a serious problem in such systems.
- Dropout is a form of regularization
- The key idea in dropout is to randomly drop neurons, including all of the connections, from the neural network during training.

# Dropout



(a) Standard Neural Network



(b) Network after Dropout

7

<sup>7</sup>Image from: [https://www.researchgate.net/figure/Dropout-neural-network-model-a-is-a-standard-neural-network-b-is-the-same-fig3\\_309206911](https://www.researchgate.net/figure/Dropout-neural-network-model-a-is-a-standard-neural-network-b-is-the-same-fig3_309206911)

# How Does Dropout Work in Practice?

- In the learning phase, we stochastically remove hidden units by setting a dropout probability for each layer in the network. We then randomly decide whether or not a neuron in a given layer is removed stochastically.

# How Does Dropout Work in Practice?

- We define a random binary mask  $m^{(l)}$  which is used to remove neurons, and note,  $m^{(l)}$  changes for each iteration of the backpropagation algorithm.



# How Does Dropout Work in Practice?

- We define a random binary mask  $m^{(l)}$  which is used to remove neurons, and note,  $m^{(l)}$  changes for each iteration of the backpropagation algorithm.
- For layers,  $l = 1$  to  $L - 1$ , for the forward pass of backpropagation, we then compute

$$a^{(l)} = \sigma(w^{(l)} a^{(l-1)} + b^{(l)}) \odot m^{(l)} \quad (1)$$

# How Does Dropout Work in Practice?

- We define a random binary mask  $m^{(l)}$  which is used to remove neurons, and note,  $m^{(l)}$  changes for each iteration of the backpropagation algorithm.
- For layers,  $l = 1$  to  $L - 1$ , for the forward pass of backpropagation, we then compute

$$a^{(l)} = \sigma(w^{(l)}a^{(l-1)} + b^{(l)}) \odot m^{(l)} \quad (1)$$

- For layer  $L$ ,

$$a^{(L)} = \sigma(w^{(L)}a^{(L-1)} + b^{(L)}) \quad (2)$$

# How Does Dropout Work in Practice?

- We define a random binary mask  $m^{(l)}$  which is used to remove neurons, and note,  $m^{(l)}$  changes for each iteration of the backpropagation algorithm.
- For layers,  $l = 1$  to  $L - 1$ , for the forward pass of backpropagation, we then compute

$$a^{(l)} = \sigma(w^{(l)}a^{(l-1)} + b^{(l)}) \odot m^{(l)} \quad (1)$$

- For layer  $L$ ,

$$a^{(L)} = \sigma(w^{(L)}a^{(L-1)} + b^{(L)}) \quad (2)$$

- For the backward pass of the backpropagation algorithm,

$$\delta^L = \Delta_a J \odot \sigma'(z^L) \quad (3)$$

# How Does Dropout Work in Practice?

- We define a random binary mask  $m^{(l)}$  which is used to remove neurons, and note,  $m^{(l)}$  changes for each iteration of the backpropagation algorithm.
- For layers,  $l = 1$  to  $L - 1$ , for the forward pass of backpropagation, we then compute

$$a^{(l)} = \sigma(w^{(l)}a^{(l-1)} + b^{(l)}) \odot m^{(l)} \quad (1)$$

- For layer  $L$ ,

$$a^{(L)} = \sigma(w^{(L)}a^{(L-1)} + b^{(L)}) \quad (2)$$

- For the backward pass of the backpropagation algorithm,

$$\delta^L = \Delta_a J \odot \sigma'(z^L) \quad (3)$$

$$\delta^l = ((w^{l+1})^T \delta^{l+1}) \odot \sigma'(z^l) \odot m^{(l)} \quad (4)$$

# Why Does Dropout Work?

- Neurons cannot co-adapt to other units (they cannot assume that all of the other units will be present)

# Why Does Dropout Work?

- Neurons cannot co-adapt to other units (they cannot assume that all of the other units will be present)
- By breaking co-adaptation, each unit will ultimately find more general features